

dablocks &
dablocks_http
~B

understanding cost of quering

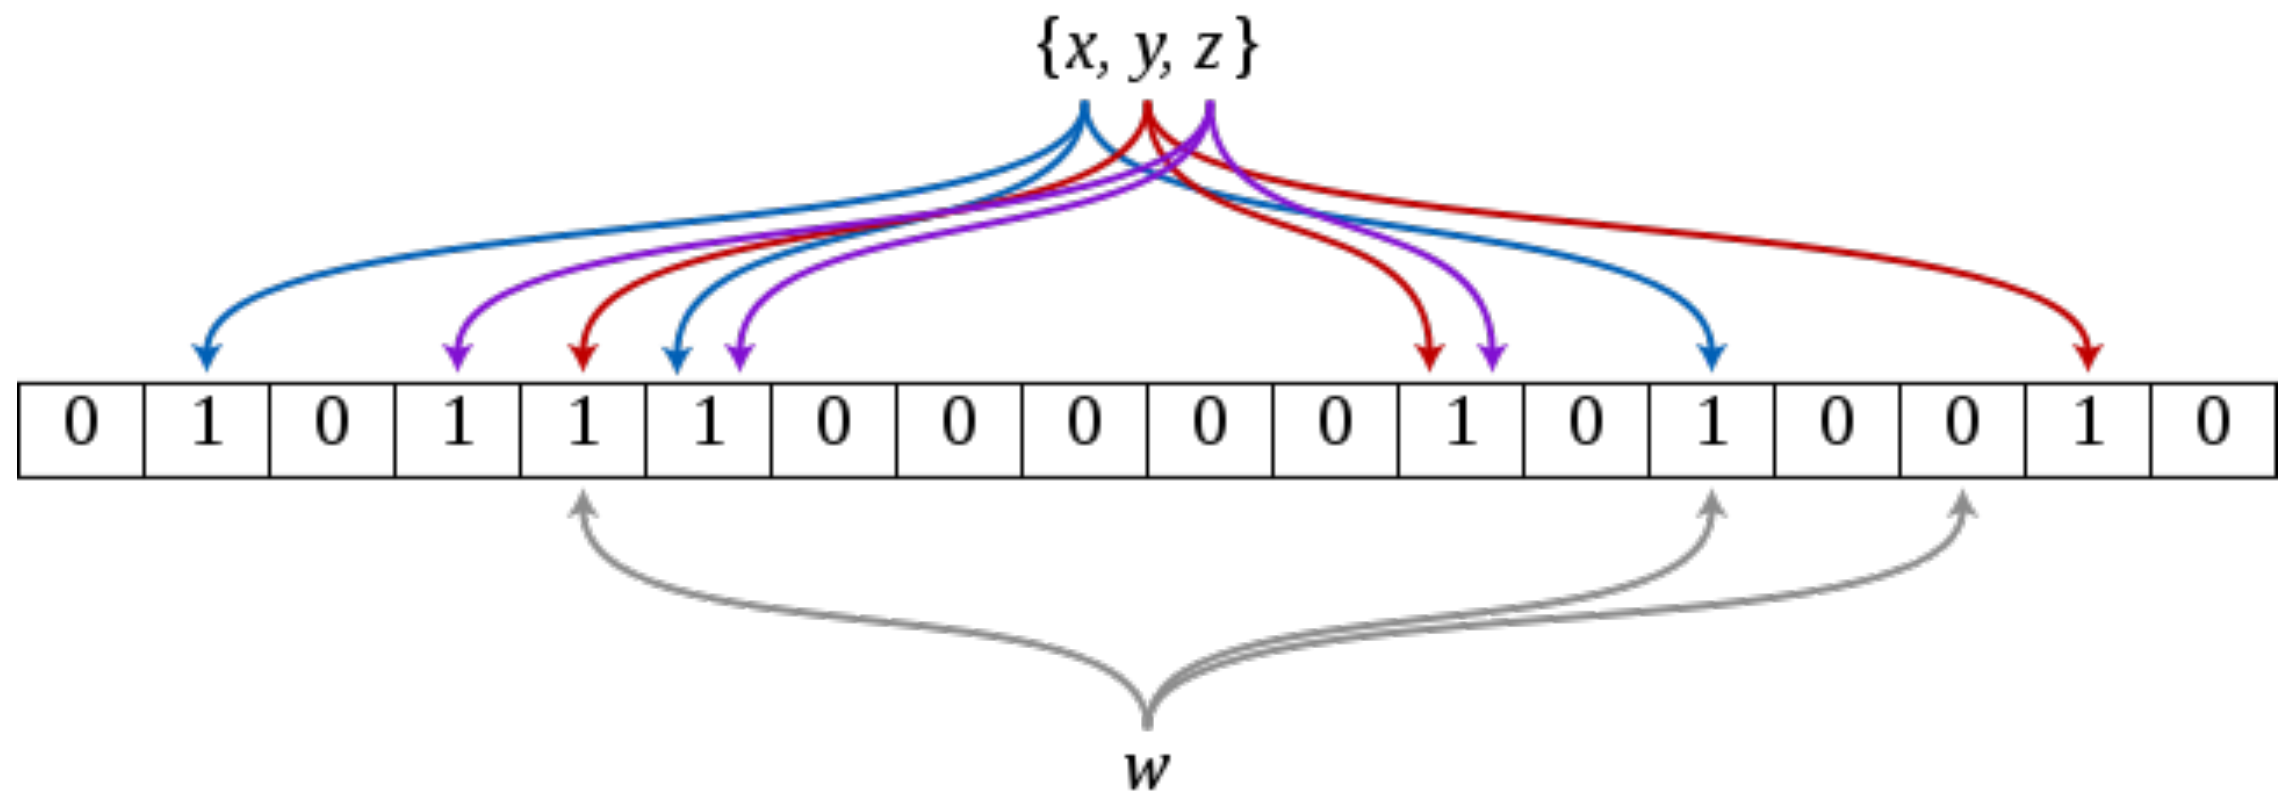
Table 2.2 Example Time Scale of System Latencies

Event	Latency	Scaled
1 CPU cycle	0.3 ns	1 s
Level 1 cache access	0.9 ns	3 s
Level 2 cache access	2.8 ns	9 s
Level 3 cache access	12.9 ns	43 s
Main memory access (DRAM, from CPU)	120 ns	6 min
Solid-state disk I/O (flash memory)	50–150 μ s	2–6 days
Rotational disk I/O	1–10 ms	1–12 months
Internet: San Francisco to New York	40 ms	4 years
Internet: San Francisco to United Kingdom	81 ms	8 years
Internet: San Francisco to Australia	183 ms	19 years
TCP packet retransmit	1–3 s	105–317 years
OS virtualization system reboot	4 s	423 years
SCSI command time-out	30 s	3 millennia
Hardware (HW) virtualization system reboot	40 s	4 millennia
Physical system reboot	5 m	32 millennia

why

- for every incoming profile id we need to check whether it exists in our DB.
- DB queries range from 20ms to 500ms or more when the DB is on another machine
- the event of creating a profile is once in the lifetime of that profile and does not change
- it's almost like flipping a switch that never turns off

bloom filter



why

- storing 100 Million profiles of keysize 2 bytes =>
=> 190 MB
- storing the same in our bloom filter with capacity for 100M entries and error rate 0.05
=> ~1.6 MB

```
[master] ~/repos/dablooms_http$ ./bin/dablooms_http /tmp/blooms/ /tmp/a.txt
2
* test scaling accuracy
Added i=100000000 ctr=100000000 words
Starting on port 9003
^C
3
[master] ~/repos/dablooms_http$ ls -lash /tmp/blooms/
total 2390288
    0 drwxr-xr-x   4 bosky  wheel   136B Apr 28 14:37 .
    0 drwxrwxrwt  12 root   wheel   408B Apr 28 14:37 ..
4 749800 -rw-----   1 bosky  wheel  366M Apr 28 14:35 bloom1.blu
1640488 -rw-r--r--   1 bosky  wheel  801M Apr 28 14:39 global.bf
```

scalable bloom filters

- Almeida, Baquero, Preguiça, Hutchison published a paper in 2006, on Scalable Bloom Filters
- suggested creating a list of bloom filters that act as one large bloom filter.
- When greater capacity is desired, a new filter is added to the list OR when a certain degree of additions also happen.

dablooms

- created by [bit.ly](#) for finding out if a url exists in their system already
- C, with wrappers in python, lisp, php, go
- instead of a bit, implements a 4-bit counter so that over time or queries, garbage collection occurs
- comes with implementation of murmur hash

dablooms_http

- C http wrapper over dablooms, created at Helpshift
- supports namespaces (each ns has its own *bf*)
- supports persisting multiple namespaces and reloading on next startup
- POST for adding entries
GET for querying

how-to

```
$ make deps compile
```

Usage:

```
./bin/dablooms_http <bloom_dir>  
<global_bloom_words_file>
```

```
$ mkdir /tmp/blooms
```

```
$ ./bin/dabloom_http  
    /tmp/blooms/ /usr/share/dict/words
```

simple membership query

```
$ curl http://localhost:9003/?key=orange
```

```
0 (NOTE: is plain/text )
```

```
$ curl -X POST -d "key=orange" http://localhost:9003/
```

```
{"ok":0} (NOTE: this is json )
```

```
$ curl http://localhost:9003/?key=orange
```

```
1
```

namespacing

```
$ curl -X POST -d "key=pune&ns=cities"  
http://localhost:9003/  
{ "ok": 1 }
```

```
$ curl  
http://localhost:9003/?key=pune  
0
```

```
$ curl  
http://localhost:9003/?key=pune&ns=cities  
1
```

examples

is this profile known?

\$ curl

http://localhost:9003/?key=profile1&ns=boomapp

have we got visits from this ip

\$ curl http://localhost:9003/?key=ip&ns=ip

is this question encountered before?

\$ curl

http://localhost:9003/?key=questionhash1&app=boom

constants.h

```
#ifndef dablooms_http_constants_h
#define dablooms_http_constants_h

#define DAEMON_ON 0
#define TEST 0
#define PORT_LISTEN "9003"

#define CAPACITY 100000000
#define ERROR_RATE .05

#define KEY_MAX_LENGTH (256)
#define KEY_PREFIX ("")
#define KEY_COUNT (1024*1024)


#endif
```

performance

after adding 10M
entries

```
#define TEST 1
```

benchmarking 30k
requests
with concurrency of
100



1

Connection Times (ms)


	min	mean[+/-sd]	median	max
Connect:	0	4 1.9	4	10
Processing:	0	4 1.6	4	10
Waiting:	0	3 1.2	2	8
Total:	3	8 1.8	8	16

Percentage of the requests served within a certain time (ms)

50%	8
66%	9
75%	9
80%	10
90%	11
95%	12
98%	12
99%	13
100%	16 (longest request)

metrics



 bosky.local.helpshift.com:9103/?metrics=1

```
{"queries":6, "hits":6, "misses":0, "additions":1, "namespaces":2}
```

references

- dablooms
<https://github.com/bitly/dablooms>
- mongoose
<https://github.com/cesanta/mongoose>
- hashmap
https://github.com/petewarden/c_hashmap
- dablooms_http
https://github.com/helpshift/dablooms_http
- Scalable Bloom Filters
<http://www.sciencedirect.com/science/article/pii/S0020019006003127>